

CITI FLAGSHIP PUBLICATIONS · AI AGENT CONSUMPTION ANALYSIS

Citi Publishes Like a Leader. AI Reads It Like a Challenger.

AI agents now read, select and relay published content before most humans see it. A first look at **how those agents consume Citi's publishing** — GPS, the Wealth Outlook, Must C and the features Flagship Publications puts into the world.

A publisher-grade read of AI consumption — not a marketing dashboard. Bespoke analysis from Innate C3, built for an editorial team, with a small live preview of Citi's own results inside.

THE OPPORTUNITY

AI Agents Are Now Your #1 Audience

Flagship Publications reaches nearly 300,000 readers. A second readership is growing faster: **AI agents** — ChatGPT, Gemini, Perplexity and the crawlers behind them — that read published content first and brief the humans who ask. Every feature Citi publishes externally is already being read, selected and quoted — or passed over — by machines. **The question for a publisher: is your work written and structured to be found by that audience?**

57.5%

of requests for web content are now bots & AI agents — machines passed humans in June 2026, 18 months ahead of forecast

1.8B

people have used AI tools worldwide

58%

have replaced search with AI for recommendations

~60%

of Google searches now end without a click

25%

projected drop in search volume by 2026

Sources: Cloudflare Radar, June 2026 (via Forbes) — Imperva's *2025 Bad Bot Report* measured its own network crossing 51% as early as 2024 · Menlo Ventures, *2025: The State of Consumer AI* · Capgemini (Jan 2025) · SparkToro, 2024 Zero-Click Study · Gartner (Feb 2024). Citi's own measured baseline follows on the next pages.

Publish for both readerships — the people, and the agents that brief them. Owned + earned, measured.

WHY THIS IS DIFFERENT

A Communications Opportunity, Not an SEO One

When an agent assembles its read on *"who produces the best research on AI, supply chains and energy security?"*, it isn't parsing keywords and meta tags. It's discovering, reading and quoting sources — three of them, and every one is editorial territory:

01

Owned publishing signals

Whether Citi's own publishing — GPS, the Wealth Outlook, Must C, the externally published features — is open, structured and quotable enough for an agent to discover and cite.

OWNED

02

Editorial authority

The outlets agents trust on markets — Bloomberg, the FT, WSJ, Reuters, Institutional Investor — and whose work they quote.

EARNED

03

Competitive context

How the sources agents trust frame the field — which houses they treat as the authorities on the economy and markets.

CONTEXT

Agents don't rank pages — they read, select and quote. Discoverability is measurable, and it's a lever the publisher controls.

WHAT WE FOUND · TWO AUDITS, ONE BRAND

One Citi, Two Reads: The Consumer Bank vs. the Publisher

We ran two small preview audits — the same method, a different lens, 10 prompts × 5 assistants each. As a **consumer bank**, Citi holds a respectable share of the AI conversation. As a **publisher of thought leadership** — home turf for Citi GPS, the Wealth Outlook and Must C — the agents are reading right past it.

	Consumer-banking lens	Thought-leadership lens
		NEW
Citi AI mindshare	38% · 19/50	16% · 8/50
Citi owned-content share	12%	1.5% · 2 of 133
Flagship series (GPS / Wealth Outlook / Must C) surfaced?	No — out of category	Yes — but rarely cited


The juxtaposition is the signal: in this sample, **the agents know Citi the bank far better than Citi the publisher** — mindshare halves, and owned share nearly disappears, precisely where Flagship Publications competes. A 50-answer preview isn't a verdict; it's a strong reason to look deeper. The rest of this deck is about that gap — and how to map and close it.

THE COMPETITIVE READ

In Our Sample, Morgan Stanley Leads the Research Answer — Citi Sits at 16%

Across 50 unbranded answers to *"who leads in economic and market research?"*, the agents reached for the same few houses. In this preview Citi sits near the bottom of the named field — a sample, not a census, but a consistent one across all five assistants.

SHARE OF VOICE · % OF 50 UNBRANDED RESEARCH ANSWERS



16%

Citi's share — trailing the leader, Morgan Stanley, by 62 points

3/10

Citi's ceiling on its best assistant — uniformly thin, no platform carries it

Below

BofA, UBS — and even Barclays & Deutsche Bank (18%)

THE EARLY SIGNAL

In this sample, Morgan Stanley, JPMorgan and Goldman anchor the AI read on research, while Citi — with GPS, the Wealth Outlook and Must C behind it — surfaces least among the major houses. **Worth testing at full scale before drawing conclusions — but the direction is hard to ignore.**


Drawn from your Innate dashboard · citation counts (Morgan Stanley 39×, J.P. Morgan 31× ... Citi 8×), shown here as share of the 50-answer sample.

OWNED CONSUMPTION

In This Sample, Agents Rarely Reached for Citi's Own Publishing

Of the **133 owned research pages** the agents cited across these answers, just **2 were Citi's** — a **1.5% owned share**. J.P. Morgan alone accounts for nearly a third. When an agent wanted a primary source to quote, it rarely pulled from Citi's shelf.

OWNED-RESEARCH CITATIONS · SHARE OF THE 133 CITED



2 of 133

Citi's research portal (citivelocity.com and peers) is cited just **twice** as an agent source. The flagship series below *do* surface — but rarely, and the agents reach for Morgan Stanley's or Goldman's publishing instead. **Citi Flagship Publications is an under-leveraged asset here, not an absent one.**

Citi GPS

Deep-dives on the structural megatrends agents are asked about most — AI, supply chains, energy security.

Wealth Outlook

The annual macro & asset-allocation outlook — exactly the “best annual outlook” query agents field.

Must C by Citi Research


Flagship market insight — world-class work the agents seldom surface when asked who leads.

SUPPORTING CONTEXT · THE BROAD LENS

The Same Pattern, Category-Wide

The publisher read is the sharper cut of a broader pattern. Measured across consumer banking, cards, travel and wealth, Citi is a solid but mid-pack name in AI — **4th at 38%**, with a **12% owned share**. The pattern is consistent across both lenses; it's simply most pronounced where Citi Flagship Publications competes.

CONSUMER-BANKING LENS · SHARE OF 50 ANSWERS



38% → 16%

The agents read Citi as a **capable everyday bank** (38%, 4th) — but as a **publisher of record**, only 16%. Same brand, same method; the more strategic the category, the wider the gap. **The consumer read is the floor; the publisher read is the opportunity.**

THE DIAGNOSIS

Our Read: A Discoverability Gap, Not a Quality One

Citi Flagship Publications stands with anyone's on substance. The pattern in this sample points to something structural — the work may not be packaged, positioned or surfaced in the form an agent can discover and quote. A fuller audit would confirm where, source by source.

WHY AGENTS MAY NOT BE REACHING IT

- Much of the best work sits **behind logins and PDFs** on citivelocity.com — invisible to the agents
- Pages aren't **structured to be quoted** — no crisp, extractable answers to the questions agents actually field
- The **franchise names** (GPS, Wealth Outlook, Must C) aren't tied to the topics agents associate with leadership

WHY THAT'S AN OPPORTUNITY

- It's a **fixable publishing problem** — the same levers that surface a Morgan Stanley note are available to Citi
- The editorial team already **owns the assets and the channel** — the publications, the newsroom, the franchises
- Progress is **measurable** — the same audit, re-run, shows the number move

Citi is doing the work. This is about making sure the agents — and the audiences they brief — can **find it, attribute it, and repeat it.**

THE OPPORTUNITY

Make Citi Flagship Publications the Source Agents Reach For

A consultative program to map the gap at full scale — 100 prompts, personas you approve — and grow from today's 16% toward the leaders, with a hard, defensible number your team can take to Citi's leadership.

16% · room to run

Position Citi Flagship Publications to be cited. Restructure and surface GPS, the Wealth Outlook, Must C and the external features as quotable, agent-discoverable answers on the exact megatrend and outlook questions agents field — so they pull Citi's publishing, not just a rival's. No fixed target; the direction is what compounds.

Measure & merchandise

A metric for the publishing investment. Each cycle produces a hard AI-citation number — proof, in a form leadership understands, that Flagship Publications is landing where credibility is now decided. A trend line to defend the budget and merchandise the wins internally.

Closing that gap is squarely the editorial team's remit — the fastest-moving lever Citi controls, with **these audits as the ongoing measure of progress.**

WHY INNATE

Software Monitors. We Interpret.

Self-service tools can tell you whether you were mentioned. They can't tell you *why* Morgan Stanley wins the research query, which of your franchises to surface first, or whether your own publishing is driving consumption — or absence.

SAAS DASHBOARDS

Monitor

- Rankings by raw mention count
- Generic, off-the-shelf prompt libraries
- A list of top sources, without strategy
- A visibility score with no "so what"

INNATE · CONSULTATIVE INTELLIGENCE

Interpret

- **Absent, present, or quoted-first?** — prominence, not just presence
- Persona-based queries designed around Citi's franchises and priorities
- Sources correlated with **wins vs. losses** — owned briefs and earned openings
- 20 years of comms judgment reading the actual answers

A FURTHER DEEP-DIVE EXAMPLE

How AI Answers the Questions Your People Are Asking

Beyond markets and products, current and prospective employees now put their questions to AI agents first. A deeper engagement can model those queries the same way — and show whether Citi's own voice, or someone else's, is shaping the answer your talent hears.

THE TALENT QUESTIONS AGENTS ALREADY ANSWER

"Is **Citi** a good place to build a career?"

"**Citi vs. J.P. Morgan** for an analyst role"

Citi's culture, pay & career growth

"Should I accept a **Citi** offer?"

Citi's hybrid-work & benefits policies

"Where is **Citi** heading under its transformation?"

WHAT A TALENT TRACK MEASURES

- Which **sources** agents lean on for these answers — Glassdoor, Reddit, Blind, the business press — and whether **Citi's own publishing and careers content** is in the mix
- How agents **frame working at Citi** vs. peer firms — and where the framing is outdated or one-sided
- Which questions have **no good Citi source at all** — the gaps your editorial calendar could fill

WHY IT'S FLAGSHIP PUBLICATIONS' OPPORTUNITY

- The features you already publish externally are **exactly the content agents could cite** for these queries — an owned answer beats a forum thread
- Measured monthly, it's an **early-warning system** for the employer narrative — and a direct input to internal comms and the exec media-training practice
- No SEO or marketing tool touches this — it's **editorial territory**

ENGAGEMENT

How We Work

Bespoke from the first step. Every prompt is modeled on how your audiences actually ask — and reviewed with you before testing.

I

Business brief
Align on priorities — the publication franchises, audiences (investors, media, clients, talent) and competitive set.

II

Prompt generation
100 audience-modeled prompts — branded and unbranded — for your review and approval.

III

Data collection
Run across your chosen 5–10 AI assistants via API; every response and citation captured in full.

IV

Analysis
Presence vs. prominence, owned vs. earned classification, competitive patterns — read by a human.

V

Presentation
Executive-ready report plus a live walkthrough with prioritized owned & earned recommendations.

100 prompts × your chosen 5–10 assistants = 500–1,000 query-runs per cycle, every one with full citation extraction.

ENGAGEMENT MODEL

Three Tiers — Take What You Need

The audit measures and recommends at a high level. The roadmap turns those findings into specific, page-level actions on Citi Flagship Publications. Execution implements them.

01	Citation Intelligence — the audit & report The measurement plus high-level recommendations . 5–10 AI assistants, 100 bespoke prompts each cycle (500–1,000 query-runs/mo); every citation classified owned vs. earned; month-over-month trends; executive briefing.	\$5-10k _{/mo}
02	Strategy Roadmap — the action plan The playbook behind the findings — a multi-page action memo (Palisociety's ran 32 pages): how to structure GPS / Wealth Outlook / Must C and the external features for agents, page-level copy & FAQ, and an outlet-by-outlet earned plan.	+\$10k _{/roadmap}
03	Execution Ongoing consultation · owned-content development — surfacing and structuring the publications for agents · technical implementation with your team · press-release & pitch development · iterative optimization.	\$15k+ _{/mo}

Audit + roadmap together: \$15–20k/mo. Citation Intelligence scales \$1k per AI assistant — 5 / 500 query-runs at \$5k, up to 10 / 1,000 at \$10k; each roadmap is +\$10k.

Modular — run one tier or all three **Wraps around** your existing agencies — proves lift, doesn't replace them **Bespoke** — no template reports

THE RETAINER IN PRACTICE

What the First Few Months Look Like

The monthly program is the point — a trend line only exists if you measure the same way over time. The measurement backbone stays constant; each month you aim it at a different question — and act on it.

EVERY CYCLE	The constant — repeatable measurement Re-run the same prompts every month, every citation classified owned vs. earned — an apples-to-apples trend line you can defend to executives.
MONTH 1	Flagship Publications sprint Surface and structure GPS, the Wealth Outlook, Must C and the external features as agent-discoverable answers on the megatrend and outlook questions where agents cite rivals today. The fastest, most-controllable lift, on the flagship franchises.
MONTH 2	Talent & employer-narrative deep dive Model the queries current and prospective employees put to agents — which sources shape the answers, whether Citi's own publishing is present, and which gaps the editorial calendar could fill.
MONTH 3	Executive & spokesperson visibility How agents portray Citi's economists and strategists vs. peer-bank voices — the researchers whose visibility is the franchise's visibility, and a direct input to the media-training practice.

A sample cadence — **you choose each month's focus**. Same retainer, a new lens; the trend line compounds.

GETTING STARTED

Let's make Flagship Publications impossible for agents to miss.

Start with the full thought-leadership audit — 100 prompts, personas you approve — to map what this preview surfaced, on the lever the editorial team controls outright. Then a new lens each month as the trend line builds.

The sample inside this deck is the 50-answer free read. The engagement is an ongoing monthly program — 100 prompts across your chosen 5–10 assistants, tailored to Citi's priorities and measured the same way every month, so the trend — and the case for the work — is real.

APPENDIX · CASE STUDY

Bain Capital Ventures: Owned Content to Displace the Listicles

BCV was absent from 82% of organic VC queries. A deep dive on every cited source found why — and the fix was an owned-content play.


84 · 69

84 "quasi-editorial" pages — directories, listicles, roundups — were among AI's most-cited sources. BCV was **absent from 69 of them** while competitors were named.

25 vs 171

BCV's own site appeared in just **25 winning answers vs a16z's 171**. The move: publish BCV's own authoritative answers to the same founder questions — a credible **owned** source for the agents to cite.

A tiered plan mapping every gap and the owned content to displace it — source-level intel a SaaS dashboard can't produce. Live dashboard → innatec3.com/bain



APPENDIX · CASE STUDY

Palisociety: Owned Content, Measured Uplift

A Marriott-affiliated boutique luxury hotel collection came to us for a baseline audit, then a roadmap that was **largely owned-media work**: exact copy changes and AI-optimized content for each property's own site — the same play Citi's publishing would run.


33% → 42%

Re-tested four months later on the newest models, AI **mention rate climbed +8 points** — top recommendations ~25% more often. One market, **Albuquerque**, jumped 35% → 75%.

Owned x2

Own-site citations doubled — the direct result of the owned-content roadmap. Earned moved too, and we could see which lever did what. Measured, attributable, repeatable.

Owned-content work, done consultatively — exact copy, per property, through the client's own review — and measured: **+8 points**. Live dashboard → innatec3.com/palisociety



APPENDIX · THE TWO SAMPLES

The 20 Prompts Behind the Two Previews

Every prompt was unbranded and category-level — no brand names — so the results reflect organic consumption, not prompted recall. A full engagement runs 100+ prompts, reviewed and approved with you before testing.

THOUGHT-LEADERSHIP LENS

- 01 Which major banks or financial institutions publish the most authoritative research on the global economy and markets?
- 02 Where do professional investors and executives turn for the best thought leadership on macroeconomic trends?
- 03 Who produces the most respected deep-dive reports on structural megatrends like AI, supply chains and energy security?
- 04 What are the best sources for institutional market analysis and cross-asset investment strategy right now?
- 05 Which institutions publish the most credible annual wealth and macroeconomic outlook reports for investors?
- 06 What research do finance professionals read to understand geopolitical and inflation risks across asset classes?
- 07 Who leads in forward-looking economic and investment thought leadership among the global banks?
- 08 Best institutional research on global capital markets and long-term investment themes this year?
- 09 Whose global perspectives and market-insights reports get cited most by the financial media and analysts?
- 10 Where can I find the most authoritative macroeconomic outlook and asset-allocation guidance for 2026?

CONSUMER-BANKING LENS

- 01 What are the best global banks for people who travel frequently and need international banking in 2026?
- 02 Which banks offer the most rewarding credit cards for everyday spending and travel right now?
- 03 How do I choose a bank that offers both personal banking and wealth management under one roof?
- 04 What do financial experts recommend for managing money across multiple countries this year?
- 05 Best credit cards for earning points on dining and travel in 2026 according to recent reviews?
- 06 Which large banks are considered most trusted for institutional financial services today?
- 07 How does wealth management at a major global bank compare to using a standalone investment firm?
- 08 What should I look for in a bank if I run a business with international clients and payments?
- 09 Most highly rated banks for combining checking accounts, savings and investment products in one place?
- 10 What are the latest innovations in consumer banking and digital financial services from top global banks?